

Hub Cultural: Mapping Cultural Event Networks in France through Instagram Graph Mining and Network Analysis

From a Diaspora Network to a General Cultural Discovery Platform

Diego Merchán

2026

Abstract

[TODO: 150–250 words. Draft last, once Results/Discussion are stable. Should state: (1) the problem — culturally-relevant activity in the Paris region (with priority given to photography, cinema, literature, theatre, and local community events) is not systematically documented or discoverable, and this project originated from, but is no longer limited to, the specific case of the Latin American diaspora; (2) the method — a pipeline that scrapes Instagram, builds a Neo4j graph, applies network analysis (PageRank, betweenness, Leiden community detection) and a three-layer NLP/LLM event-extraction pipeline, combined with human curation as the trust signal, plus a research-operations discipline (decision log, run log, dry-run-first design) treated as methodology in its own right; (3) headline findings — network size and structure, the bridge role of institutional accounts, event landscape, and a transparent account of data-quality limits; (4) the practical artifact — a public, bilingual event-discovery website; (5) contribution — methodological (human curation vs. automated scoring; a cultural-relevance criterion replacing a nationality-based one) and empirical (first mapped network of this kind for this cultural ecosystem in France). **]**

Contents

1	Introduction	3
1.1	Context and Motivation	3
1.2	Problem Statement and Research Question	4
1.3	Objectives	4
1.4	Scope and Contributions	5
1.5	Thesis Structure	6
2	Literature Review / Theoretical Framework	7
2.1	Diaspora Studies and Cultural Networks	7
2.2	Cultural-Sector and Creative-Industry Network Mapping	7
2.3	Social Network Analysis in Migration and Cultural Contexts	8
2.4	Computational Social Science and Social Media as Data	8
2.5	Graph Databases and Network Science Methods	8
2.6	NLP for Event Extraction and Classification	9
2.7	Positioning	9
3	Methodology	10
3.1	Research Design Overview	10
3.2	Research Operations: Decision Logging, Run Tracking, and Iterative Safety	11
3.3	Data Collection	13
3.4	Data Model	14
3.5	Network Analysis Methods	14
3.6	Event Detection Pipeline	15
3.7	Event Deduplication	16
3.8	Evaluation Design: Blind Human-versus-Pipeline Comparison	17
3.9	Human Curation as the Trust Signal	19
3.10	Geospatial Enrichment	20
3.11	Ranking and Relevance Scoring	20
3.12	Ethical Considerations	21
3.13	Limitations of Method	22

4	System Implementation: The Hub Cultural Platform	23
4.1	Pipeline Architecture Overview	23
4.2	Extraction and Ingestion	23
4.3	Network Analysis Engine	24
4.4	Event Enrichment Services	24
4.5	The Discovery Website	24
5	Results	26
5.1	Network Structure	26
5.2	Bridge Actors and Community Structure	27
5.3	Thematic and Event Landscape	28
5.4	Geographic Distribution	29
5.5	Data Quality Findings	29
6	Discussion	31
6.1	Interpreting Network Structure	31
6.2	From a Diaspora Network to a Cultural-Relevance Network	31
6.3	On Data Quality and the Limits of Automation	33
6.4	Practical Implications: The Website as Intervention	34
6.5	Methodological Reflections	34
7	Conclusion	35
7.1	Summary of Contributions	35
7.2	Answering the Research Question	35
7.3	Limitations	35
7.4	Future Work	35
7.5	Closing Reflection	36
A	Graph Schema Reference	37
B	Selected Decision Log Entries	38
C	Evaluation Protocol and Samples	39
D	Discovery Website — Screenshots	40
E	Repository Structure	41

Chapter 1

Introduction

1.1 Context and Motivation

Hub Cultural began as an attempt to make visible something that is well known anecdotally but poorly documented empirically: the cultural life that Latin American, and in particular Colombian, communities sustain in the Paris region through small associations, independent venues, cultural attachés, and a dense but scattered presence on Instagram. The author’s own Colombian background supplied both the initial motivation and the first seed account for data collection — the Colombian consulate in Paris (@consuladocolparis) — from which the network was grown outward through automatic relationship discovery.

As the project progressed, the boundary that had organized it from the start — *is this account or event part of the Latin American diaspora?* — proved both practically brittle and conceptually narrower than what the data, and the underlying research interest, actually supported. Many of the best-connected cultural accounts discovered through the network turned out to be French institutions, mixed-nationality collectives, or venues with no diaspora affiliation at all, but which were unmistakably part of the same cultural ecosystem the project set out to map — programming photography exhibitions, independent cinema, literary events, and theatre alongside explicitly Colombian or Latin American programming. The project’s scope therefore evolved in two steps: first from Colombia-specific to Latin-American-diaspora-wide, and subsequently from a nationality-based criterion to a purely cultural one. The operative question became not *who is behind this account* but *is this culturally relevant*, with explicit priority — inherited from the project’s original four thematic agendas (docs/objetivos.md) — given to photography, cinema, literature, and theatre, alongside a fifth, less formally bounded category: events organized by and for local communities, regardless of national or cultural origin.

This evolution is treated in this thesis not as scope creep to be apologized for, but as a methodological finding in its own right (§6.2): a nationality-based inclusion criterion

turned out to be both harder to operationalize consistently and less faithful to how cultural networks actually cohere in practice, where diaspora and host-society cultural production overlap and reinforce each other rather than remaining cleanly separable. It also turned out to be measurably encoded in the pipeline itself: the semantic reference set used by the first stage of event detection was written entirely in the vocabulary of Colombian diaspora events, and a blind evaluation of 100 posts showed that non-diaspora cultural accounts consequently lost real events at roughly three times the rate of diaspora accounts (§3.8). The scope question was therefore never purely conceptual; it had a direct, quantifiable effect on what the system was able to see.

1.2 Problem Statement and Research Question

The cultural activity produced by, and adjacent to, migrant and diaspora communities in a host country is rarely centralized: it lives across dozens of Instagram accounts, is announced inconsistently, and is legible mainly to whoever already follows the right accounts. This is true of Latin American cultural life in the Paris region specifically, but — as this project’s own data eventually made clear — it is not a problem unique to any one national community; it is a structural feature of how small-scale cultural production circulates on social media generally.

This thesis asks: *How is culturally-relevant activity — with priority given to photography, cinema, literature, theatre, and local community cultural life — organized as a network of accounts and events in the Paris region, and can that structure be systematically mapped, analyzed, and made discoverable using publicly available Instagram data?*

This breaks into three sub-questions:

1. **Structural.** Who are the connective actors in this network — which accounts function as hubs or bridges between otherwise separate cultural circles, and does this differ between institutional accounts and individual or grassroots ones?
2. **Thematic.** How is cultural activity distributed across the project’s priority domains (photography, cinema, literature, theatre, community events), and how well does an automated pipeline classify it into these domains compared to human judgment?
3. **Practical.** Can the resulting graph be turned into a public-facing tool that helps the community it describes actually discover and attend these events — and what does building that tool reveal about the reliability of the underlying data?

1.3 Objectives

The project’s objectives, as originally framed (`docs/objetivos.md`) and updated here to reflect the current, cultural-relevance-based scope:

1. Build and maintain four thematic monitoring agendas — *literature*, *photography*, *cinema*, and *theatre* — tracking the organizations that play an active role in each, complemented by a fifth, less formally bounded priority: cultural events organized by and for local communities.
2. Consolidate a pipeline that extracts events from the accounts under these agendas, primarily from Instagram and, prospectively, from websites and other social networks.
3. Map the network of actors within each thematic agenda and categorize them — by type of organization, geographic zone, cultural identity where relevant, and structural role in the network.
4. Evaluate the participation opportunities these networks open up, using a deliberately human, non-technical success criterion: on this dimension, the project succeeds when it leads to actually meeting someone through it.

1.4 Scope and Contributions

Scope. Geographically the project centers on Île-de-France, though a meaningful share of curated accounts and events originate from, or are organized by, institutions physically located outside France — a scope boundary discussed explicitly in §3.13 rather than silently discarded. Instagram is the sole data source, a deliberate boundary rather than an oversight: it is the platform where this cultural activity is, in practice, announced. Only public accounts are in scope. Inclusion is governed by cultural relevance rather than national origin, with the four thematic pillars and the community-events category acting as explicit priorities rather than hard filters.

Contributions. This thesis makes three kinds of contribution:

- *Methodological* — a working pipeline combining network analysis, multilingual NLP event extraction, and human curation as the operative trust signal, arrived at through an evidenced pivot away from fully automated account scoring (§3.9); a blind human-versus-pipeline evaluation protocol used to drive successive redesigns of the extraction logic (§3.8); and a research-operations practice — a running decision log, a run log, and a dry-run-first discipline applied to every script — treated here as part of the methodology rather than incidental bookkeeping (§3.2).
- *Empirical* — a systematic mapping of this cultural network’s structure: node and edge counts, community structure, and the role of bridging actors, alongside a transparent account of where the underlying data is, and is not, reliable (geocoding confidence, date-extraction accuracy, categorization accuracy).

- *Applied* — a public-facing, bilingual event-discovery website that turns the graph into something a member of the community it describes can actually use, not only an analytical artifact for the researcher.

1.5 Thesis Structure

Chapter 2 situates the project within diaspora and migration studies, cultural-sector network mapping, social network analysis, computational social science, and NLP-based event extraction. Chapter 3 details the methodology, including the three methodological pivots this thesis treats as contributions in their own right: human curation as the trust signal, blind human-versus-pipeline evaluation as the driver of redesign, and a research-operations discipline of decision logging, run logging, and dry-run-first iteration. Chapter 4 describes the resulting system, from data collection through the public discovery website. Chapter 5 reports results: network structure, thematic and geographic distribution, and a transparent accounting of data quality. Chapter 6 discusses what these results mean, including a dedicated reflection on the project’s evolution from a nationality-based to a cultural-relevance-based scope. Chapter 7 concludes, states limitations, and outlines future work.

Chapter 2

Literature Review / Theoretical Framework

This chapter frames the project along two axes that the project's own history forces it to hold together. The first is the diaspora-studies axis, which motivated the work and still describes a substantial part of the corpus. The second is the cultural-sector axis: since the operative inclusion criterion is now cultural relevance rather than national origin (§6.2), the theoretical framing has to account for the network as a field of cultural production in a metropolitan area, not only as a diaspora formation. The remaining sections supply the methodological grounding — network analysis, social media as data, graph databases, and NLP event extraction — that Chapters 3 and 4 draw on.

2.1 Diaspora Studies and Cultural Networks

[TODO: Situate within diaspora/migration studies: how migrant communities organize culturally, the role of associations and informal networks, prior work on Latin American diasporas in Europe specifically if available. One concrete hook already present in the project's own notes: DD-020 in `docs/decisions_es.md` reads the graph's low density as potentially sociological rather than purely technical — a very sparse network as a possible sign of diaspora fragmentation in a host country with high individualisation, citing Vertovec (2009). That citation needs to be checked and properly sourced before use, and is currently the only literature reference recorded anywhere in the repo. **]**

2.2 Cultural-Sector and Creative-Industry Network Mapping

[TODO: New section (was folded into the diaspora section in the first outline). Needed because the project's inclusion criterion is now cultural relevance, not national origin:

cover prior work mapping cultural/creative-sector actors as networks (cultural intermediaries, venue/institution ecologies, scene and field theory, cultural participation and discoverability). This is the literature that justifies treating a heterogeneous set of venues, institutions, collectives and individuals as one object of study — and against which the thesis’s empirical claim (a small institutional core with a large organic periphery, §5.2) should be read.]

2.3 Social Network Analysis in Migration and Cultural Contexts

[**TODO:** Prior applications of SNA to migrant/diaspora networks and to cultural fields; concepts to define here for later use: degree, betweenness and PageRank centrality, community detection and modularity, structural holes and brokerage, participation coefficient, and the E-I index. All five are computed in `3_analyze_network.py` and reported in §5.1, so each needs a definition here that the results chapter can rely on without re-explaining.]

2.4 Computational Social Science and Social Media as Data

[**TODO:** Justify Instagram as a data source; discuss known biases (who is and is not represented on Instagram, algorithmic visibility, platform terms of service, and the ethics of public-data research). Two project-specific points to connect to the literature: (i) the corpus mixes a behavioural signal (mentions, tags, co-authorship) with a platform-inferred one (Instagram’s own `relatedProfiles` suggestions), and this thesis keeps them as separate network layers rather than merging them (§3.5); (ii) scraping is mediated by a commercial actor (Apify), which is itself a methodological dependency worth naming.]

2.5 Graph Databases and Network Science Methods

[**TODO:** Neo4j and the property-graph model; why a graph database rather than a relational one for this problem (DD-001); brief technical grounding, with original references, for PageRank, betweenness centrality, the Leiden algorithm, k-core decomposition, and the participation coefficient.]

2.6 NLP for Event Extraction and Classification

[TODO: Ground the three-layer detection design: sentence embeddings for candidate filtering, zero-shot NLI classification, and LLM-based extraction, classification and gating. Cite relevant NLP literature (sentence embeddings and semantic similarity, zero-shot NLI, LLM information extraction, and the reliability problems of LLM output), and treat the multilingual requirement (ES/FR/EN) as a specific design constraint rather than an implementation detail — DD-008 records a monolingual English NLI model producing a 0.7% detection rate on Spanish and French captions, which is the empirical anchor for this argument. **]**

2.7 Positioning

[TODO: One or two paragraphs. Three differentiators are already supported by the material in this thesis: (i) against prior diaspora network studies — a data-driven, automated-plus-human-curated hybrid that produces a live public tool rather than only an analysis; (ii) against prior event-extraction NLP work — domain-specific gating logic validated by blind human comparison at three successive scales (§3.8), and the finding that the semantic reference set encoded a cultural bias measurable in recall; (iii) against both — the explicit replacement of a nationality-based inclusion criterion by a cultural-relevance one, including the removal of a nationality bonus from the ranking formula (§6.2). **]**

Chapter 3

Methodology

3.1 Research Design Overview

The project is organised as a seven-phase pipeline, each phase implemented as one or more independent, re-runnable scripts writing to a shared Neo4j graph:

1. **Extraction** — Instagram profiles and posts are scraped through Apify actors into local JSON (`1_harvest_ig_profiles.py`, `1_harvest_ig_posts.py`, `extract_profiles.py`).
2. **Ingestion** — the raw JSON is loaded into Neo4j Aura (`2_build_graph.py`).
3. **Network analysis** — the graph is exported to CSV and analysed offline with `igraph` and `leidenalg`, and the resulting metrics are written back (`3_analyze_network.py`).
4. **Manual categorisation** — a hand-curated spreadsheet of account attributes is uploaded onto the matching `:Account` nodes (`load_manual_account_categorization.py`).
5. **Event enrichment** — posts are passed through the three-layer detection pipeline to create `:Event` nodes, with inline deduplication (`4_enrich_events_extract.py`), followed by a standalone deduplication pass over existing events (`4_enrich_events_resolve.py`).
6. **Geographic enrichment** — location strings are geocoded and placed in an administrative hierarchy (`4_enrich_locations.py`).
7. **Cleanup** — accounts never touched by manual curation, and everything depending exclusively on them, are deleted (`cleanup_legacy_accounts.py`).

Two design principles run through all seven phases and are treated in this thesis as methodological rather than merely technical. *Idempotency*: every script can be re-run

without duplicating or corrupting what previous runs produced, because each checks a sentinel property before processing (`eventExtracted`, `lat`) or merges rather than inserts. *Separation of provenance*: the manual curation phase deliberately writes to a property namespace (`artType`, `culturalIdentity`, `geoZone`, `manualFollowersCount`) disjoint from the one the ingestion phase writes, so that neither can silently overwrite the other. The practical consequence is that automated and human-produced data remain distinguishable in the graph after the fact, which is what makes the curation-as-trust-signal argument of §3.9 verifiable rather than merely asserted.

[TODO: Insert the pipeline diagram here, adapted from `docs/data_flow.md` (an English version already exists alongside `docs/data_flow_es.md`). **]**

3.2 Research Operations: Decision Logging, Run Tracking, and Iterative Safety

Independently of the technical pipeline itself, this project maintained three interlocking practices for managing the research process, treated here as a methodological contribution rather than incidental project management. Together they make the pipeline’s evolution auditable after the fact — which is not free in a setting where the object of study is a continuously growing graph, reprocessed repeatedly by successive versions of the extraction logic.

A real-time decision log. Every non-trivial design choice was recorded, at the moment it was made, in a sequentially numbered decision log (`docs/decisions_es.md`, forty-one entries numbered DD-001 to DD-044; three numbers in that range were reserved for decisions referenced in the code but never written up) rather than reconstructed retrospectively for this thesis.¹ Each entry states the alternatives considered, the evidence that motivated the final choice, and, where applicable, validation results. Four entries illustrate the range of what this captured. DD-041 documents a geographic-conflict guardrail added to the event-deduplication resolver after a production dry-run surfaced a specific failure mode: two structurally similar events merged despite being in different countries. It was validated first against fourteen isolated test cases and then in a live dry-run over the 664 events then in the graph, in which it blocked six geographically contradictory merges while leaving intact the 68 duplicate pairs the run had identified for merging. DD-042 documents the addition of a per-event, LLM-generated tag field, designed specifically to avoid a parsing failure present in an earlier, account-level free-text field, and scoped explicitly as prospective — the events already in the graph were left without it rather than backfilled at token cost. DD-043 and DD-044 document the design of the

¹The log is maintained in Spanish, the language in which the project work was conducted; `docs/decisions.md` holds a partial English translation that currently stops at DD-033. Where entries are quoted in this thesis, the translation is the author’s.

Neo4j-to-static-JSON export and the resulting public website, including an explicit accounting of which ranking components could be precomputed at export time and which had to remain client-side because they depend on the current date or on the individual visitor’s session. Maintaining this log in real time, rather than writing it up after the fact, is itself a methodological choice: it captures the reasoning available *at the time* of each decision, including alternatives that were considered and rejected — something retrospective documentation systematically loses.

A run log. Separately, every substantive execution of the pipeline against live data — not every script invocation, but every run whose output was retained or acted upon — was logged with its date, parameters, and a summary of results (`docs/runs_log_es.md`, RUN-001 to RUN-021 at the time of writing). This served two purposes beyond record-keeping. It made it possible to distinguish a genuine change in the underlying data (new accounts or events discovered) from a change in pipeline logic when comparing two exports — a distinction that becomes impossible to reconstruct once several months of both have accumulated. And it created a paper trail for decisions that depended on a specific run’s output rather than on memory: the decision not to add a venue-specificity guardrail to the event resolver, for instance, is recorded in RUN-021 with the concrete case that prompted it (one event node absorbing three separate watch parties for the same football match, held in different Paris bars) and the reasoning for accepting it.

Dry-run-first, diagnose-first design. Every script that writes to the live Neo4j graph was built with a `-dry-run` or `-diagnose` mode as a first-class feature, not an afterthought: `cleanup_legacy_accounts.py` runs its deletion logic inside a single transaction and issues a `ROLLBACK` under `-dry-run`, so a preview’s reported counts are exact rather than estimated; `load_manual_account_categorization.py` and the event-extraction and event-resolution scripts all follow the same pattern. Combined with the idempotent design of the extraction and ingestion scripts, this made it possible to iterate on pipeline logic against a live, shared, non-trivially-sized research database without the usual risk of an irreversible mistake. The discipline was not self-enforcing, however, and DD-038 records what happens when it lapses: the diagnostic CSV export was gated on `-dry-run`, so the first real production run of the event extractor produced no comparison file and — because event timestamps were written as strings rather than native Neo4j datetimes — could not be reliably queried by date afterwards either. A run that worked correctly left almost no auditable trace. The fix restored parity between the two modes, on the principle that a production run needs to be at least as inspectable as a rehearsal.

Together, these three practices form a coherent research-operations discipline: decisions are recorded when made, runs are recorded when executed, and changes to the live data are always previewable before being committed. For a thesis built on a continuously evolving dataset rather than a fixed one, this discipline is what makes it possible to state, for any given result in Chapter 5, exactly which version of the pipeline and which run

produced it.

3.3 Data Collection

Data collection relies on Apify's hosted `instagram-profile-scraper` and `instagram-post-scraper` actors rather than on a self-maintained browser-automation stack (DD-002). The reasoning recorded at the time was one of maintenance cost and risk: Instagram actively blocks direct scraping, so a self-hosted scraper implies continuous maintenance and IP-blocking risk, whereas a hosted actor externalises both at a metered per-use price. That price is not incidental to the method — it is tracked per run in `.apify_cost_log.json` and functions as a real constraint on sampling decisions, which is why the harvesting logic is built around spending as little as possible per unit of new information. Profile scraping and post scraping are separate scripts (DD-003) because they use different actors, return different JSON structures, and need different refresh frequencies.

Seed selection changed across the project's two phases. The first phase used a single seed, the Colombian consulate in Paris, and expanded outward breadth-first through mentions, tags and Instagram's own `relatedProfiles` suggestions (DD-006); the reasoning was that a consulate is the most central institutional node of the community it serves, so relevant accounts should lie within one to three hops, and that organic expansion avoids the researcher-bias of a hand-written list. The second phase replaced the single seed with curated seed files (`config/seeds_*.json`) covering Latin American consulates and embassies in France plus French cultural and academic institutions connected to Latin America (DD-022). These institutional seeds raised a conflict the log resolves explicitly: the account-classification scheme treated government accounts as out of cultural scope, yet the entire discovery process depends on them. DD-026 separates the two questions — *is this account a valid target for cultural analysis?* and *is this account structurally necessary to discover the network?* — and assigns consulates the role of seed source without treating them as cultural targets.

Post collection is deliberately recency-biased rather than density-maximising. Rather than retrieving a fixed number of historical posts per account, the harvester computes a per-account time window equal to the gap since that account's most recent known post, capped at a configurable maximum, and merges results into a rolling window of the fifty most recent posts (DD-028, DD-029). The rationale is that the downstream objective is detecting forthcoming events, not reconstructing publication history; the earlier fixed-window design had a blind spot in which an account with a recent post was skipped entirely, losing anything published in between. A related robustness fix is worth recording as method rather than trivia: the Apify actor signals "no content found" by returning a one-element list containing an `error` key, which is `truthy` in Python and was therefore

silently treated as a successful single-post result, corrupting both the run counters and — before the merge logic existed — the stored file (DD-030). Data collection failures in this setting do not always announce themselves as failures.

3.4 Data Model

The graph uses eleven node labels and sixteen relationship types. Accounts, posts, hash-tags, locations, music tracks, comments and IGTV videos come from ingestion; events come from the NLP pipeline; cities, arrondissements and countries come from geocoding. Accounts additionally carry `:Public` or `:Private` as a second label rather than a boolean property (DD-005), which allows indexed filtering in Cypher and visual separation in the Neo4j browser. The full schema is reproduced in Appendix A.

The design principle behind the node set (DD-004) is that the model should capture three distinct dimensions at once: content (posts, hashtags), social interaction (mentions, tags, co-authorship), and cultural context (locations, music tracks). A minimal account-and-post model was considered and rejected on the grounds that it would discard exactly the signals the research question depends on — venues, collaborations and thematic markers. One consequence of this model matters for the analysis chapter and is easy to overlook: social relationships such as `MENTIONS` and `TAGS_USER` originate at the `:Post` node, not at the account, so any account-to-account social network has to be obtained by projecting `author→post→target` rather than read directly off the stored edges (§3.5).

3.5 Network Analysis Methods

Network metrics are computed offline. The graph is exported to CSV, the metrics are computed locally with `igraph` and `leidenalg`, and the results are written back to Neo4j in batches. This is a deviation from the conventional approach of running Neo4j's Graph Data Science library in-database, and the reasons are recorded in DD-013: the Bolt port (7687) that GDS requires is blocked on the network the work was carried out from, standard AuraDB does not include the GDS plugin at all, and — the argument that made the constraint into a preference rather than a workaround — CSV outputs are version-controllable and re-analysable without any database connection or cloud credentials. The earlier GDS-based script is retained in `old/` for reference but never ran successfully in this project's actual environment, which also means the "Cultural Relevance Score" it computed is a documented design, not a produced result.

The network is treated as **multiplex**: two layers analysed separately rather than merged. The *social layer* contains organic interactions — mentions, tagged users and co-authorship — projected `author→post→target` as described above. The *algorithmic layer* contains Instagram's own `relatedProfiles` similarity suggestions (`RELATED_TO`).

Keeping them apart is a substantive methodological choice rather than a technical convenience: the social layer records who actually engages with whom, while the algorithmic layer records what a proprietary recommender considers similar. Merging them would blend a behavioural signal with a platform inference, and any finding about brokerage or community structure would then be partly a finding about Instagram’s recommendation engine.

On each layer the analysis computes exact PageRank and exact betweenness centrality (not approximations, which the graph size permits), Leiden community detection at three resolutions ($\gamma = 0.5, 1.0, 1.5$ with a fixed seed of 42 for reproducibility), weakly connected components, k-core decomposition, the participation coefficient as a brokerage measure, and the E-I index of internal versus external mixing, computed both by curated actor type and by detected community. Running Leiden at three resolutions rather than one is a hedge against resolution-dependent artefacts: a community structure that survives all three is a stronger claim than one tuned to a single parameter. Centrality is computed over the whole graph rather than over a filtered subgraph, with tier filters applied only at the reporting stage (DD-014), because an account’s PageRank depends on who mentions it, including accounts that are themselves out of scope; filtering first would measure the importance of accounts within a hand-picked subset rather than within the network they actually inhabit.

3.6 Event Detection Pipeline

Event detection runs in three layers over post captions. The design was originally four layers and was reduced to three in August 2026 after the evaluation described in §3.8.

Layer 1 — semantic candidate filtering. Each caption is embedded with `paraphrase-multilingual-MiniLM-L12-v2` and compared by cosine similarity against a reference set of roughly 100 phrases covering several event types, three languages and six announcement forms (calls, openings, listings, date-plus-place, registration, reminders). Embeddings rather than keywords were chosen because the corpus mixes Spanish, French and English with emoji and slang, where keyword matching fails on orthographic variation and indirect phrasing (DD-007); the reference set is large rather than minimal (DD-010) because a narrow reference space misses exactly the small community events the project most wants to capture. Layer 1 exists primarily as a cost control: running the NLI model over every post was measured at five to seven hours on the available hardware, whereas pre-filtering reduces the candidate set to roughly 40% of the corpus and the total to about eighteen minutes (DD-009).

Layer 2 — is this event-like at all. Surviving candidates are passed to a binary natural-language-inference classifier (`MoritzLaurer/multilingual-MiniLMv2-L6-mnli-xnli`). Two decisions here are

worth stating explicitly. The model is multilingual by requirement, not by preference: an earlier English-trained NLI model produced a detection rate of 0.7% on Spanish and French captions — effectively noise (DD-008). And the hypothesis sentence is now formulated in the caption’s own detected language, after a fixed-Spanish hypothesis was shown to depress scores for English and French captions regardless of their content.

Layer 3 — LLM extraction, typing and gating. Candidates that pass Layer 2 are sent to a large language model (Ollama locally, or Groq or Cerebras in the cloud, with automatic failover between the two cloud providers and a cap on rate-limit waiting, DD-033) which extracts title, description, date, location and price range, assigns an event type from a fixed sixteen-label taxonomy, proposes one to three free-form artistic tags (DD-042), and — the decisive change — *gates* node creation. If the model reports that the post is not a public invitation, or not upcoming, no `:Event` node is created at all. In the earlier design the node was created regardless and merely scored lower, which §3.8 shows was the single largest source of false positives.

Three further constraints shape the layer-3 stage. Calls are deduplicated by post identifier within a run, so co-published posts share one call. Reasoning traces are requested only in dry-run mode, since they are expensive and only useful for auditing. And a deterministic `has_text_date` gate — requiring that a date actually appear in the caption text — is evaluated *before* the model is called rather than after, which was measured to remove 40.7% of calls offline and 48.3% in production without changing a single final decision (DD-036). The date extractor behind that gate is itself constrained to avoid manufacturing precision: a month or year mentioned without a day is discarded rather than resolved to an arbitrary first-of-month (DD-034), after four disagreements in the evaluation sample were traced to exactly that behaviour.

Newly created events inherit `artType`, `institutionType`, `culturalIdentity`, `geoZone` and `parentInstitution` directly from the publishing account’s manual curation, at no model cost, whenever that account has been curated. Deduplication happens twice: an inline resolver checks cosine similarity above 0.82 together with a location match and a date within three days before creating a new node rather than enriching an existing one, and a standalone resolver (§3.7) runs afterwards over the accumulated set.

3.7 Event Deduplication

Several accounts routinely publish about the same event, so deduplication is not an edge case but a structural feature of the corpus. The original resolver grouped events by normalised location string and compared pairs only within a group, using date proximity and cosine similarity as secondary criteria (DD-011). An audit of a production run showed this ordering to be the wrong way round: two posts about the same concert written the venue address differently and therefore never met, and an event with no location string

at all could never be compared with anything (DD-040).

The redesign inverts the order. All pairs are compared by embedding similarity first, using a vectorised cosine matrix, with dates as mandatory confirmation rather than as a primary filter: where both dates exist they must fall within three days and similarity must exceed 0.75; where one is missing, a stricter similarity threshold of 0.85 compensates for the lost corroboration. Location was explicitly *not* converted into an embedding, on the reasoning that text embeddings capture lexical rather than geographic similarity — two different theatres sharing the words "Théâtre de" may embed close together, while two addresses on the same street may not. An audit of the location strings then in the corpus supported demoting them: roughly a fifth carried a full address, over half only a venue name, a tenth only a city or country, and around a seventh were empty.

The predictable cost of dropping the location filter appeared on the first live dry-run: merges that were geographically contradictory, such as a Paris event merging with a Madrid one, because generic community posts share vocabulary regardless of where they take place. The response (DD-041) was a deliberately small gazetteer of country and city names drawn from the actual failure cases, which blocks a merge only when both sides name a recognisable and non-overlapping country. It is permissive by design — insufficient evidence on either side never blocks — precisely to avoid reintroducing the exact-match problem the redesign removed. Its residual limit is documented rather than patched: it compares countries, so two different venues within the same city can still merge, which RUN-021 shows happening and which the run log records as an accepted trade-off for the case at hand.

3.8 Evaluation Design: Blind Human-versus-Pipeline Comparison

The pipeline's accuracy was assessed by blind manual classification of post captions against the pipeline's own decisions, repeated at three increasing scales, with each round feeding concrete changes back into the code. In each round the captions were classified without access to the pipeline's decision columns, and only then joined against the diagnostic CSV the extractor emits.

Round 1 — 100 posts (2026-08-09). Overall agreement was 58%. The asymmetry mattered more than the headline: 35 disagreements were posts the pipeline called events and the human classification did not, against 7 in the opposite direction. Of those 35 false positives, 29 already carried the layer-3 model's own verdict that the post was neither a public invitation nor upcoming — the system had the information and did not act on it, because node creation was not conditioned on those flags, only scoring was. Simulating the gate on the same sample raised precision from 56% to 90% and overall agreement

from 58% to 86%. This single finding is what turned gating from a scoring adjustment into a creation condition (§3.6).

The same round produced the bias finding referred to in §1.1. The reference phrases of layer 1 had been written entirely in the vocabulary of Colombian diaspora events. Splitting the sample by account theme, the false-negative rate was 4.3% for Colombian and Latin American accounts against 12.9% for general and French accounts — real events with explicit dates and venues, in French and English, discarded just below the similarity threshold. A criterion that was intended to be about cultural relevance was, at the level of implementation, partly about national vocabulary.

Round 2 — 495 posts (batches through 2026-08-12). Aggregate precision was 0.824 and recall 0.858. Three code changes followed (DD-036): eighteen additional reference phrases in an institutional and formal register, added after the three layer-1 false negatives in the sample were found to share a sober institutional tone absent from the reference space; the deterministic `has_text_date` gate, which was shown offline to block 15 of the 27 false positives in the full sample; and an explicit negative instruction in the layer-3 prompt excluding course registrations, enrolment and admission calls, which had been producing events with a structure indistinguishable from real ones. A fourth change was considered and rejected on the evidence: lowering the layer-1 threshold from 0.45 to 0.35 would have recovered seven real events at the cost of admitting 47 additional non-events, a signal-to-noise ratio of roughly 1:6.7, which is the wrong trade in a pipeline whose dominant error was already precision.

Round 3 — 150 posts in a live dry-run (2026-08-12, RUN-020). Agreement reached 96.7%, with precision 97.3% and recall approximately 90% — the first round run against the modified pipeline rather than a simulation. Four of the five remaining disagreements were traced to specific lines of code and fixed the same day (DD-037): a theatre season written "26/27" parsed as a day-and-month date, and a 600-character truncation window in the date extractor that discarded dates appearing later in the long multi-line captions typical of event flyers. The fifth had no identifiable code cause — two nearly identical posts about the same event, one accepted and one rejected by the language model — and is recorded as observed variability rather than resolved.

Two methodological caveats belong with these numbers. First, the human labels are the author's own, so this measures agreement with a single annotator's criterion, not with a consensus ground truth; the criterion itself had to be adjudicated in at least one case, where exhibitions announced with an explicit closing date were ruled to count as events even when the surrounding text is retrospective (DD-035), and that adjudication remains flagged as unresolved against the pipeline's own behaviour. Second, agreement rose across the three rounds partly because the pipeline improved and partly because the samples differ in composition, so the figures track a trajectory rather than three measurements of the same quantity.

3.9 Human Curation as the Trust Signal

The project's original design selected accounts automatically. A classifier scored each discovered account on geographic and cultural dimensions, combining semantic similarity against reference embeddings with hard signals from the profile, and accounts above a threshold were kept. In practice, well under 10% of the accounts it selected turned out to be culturally valuable on manual review. The response was not to tune the classifier further but to change what the pipeline treats as authoritative: the property `manualDataCuratedAt`, written only by `load_manual_account_categorization.py` when an account appears in the hand-curated spreadsheet, became the operative signal for "this account is worth keeping", and the automated score was retired from that role.

This is a pivot worth stating as a finding rather than as a confession. Three mechanisms visible in the project's own records explain why the automated criterion underperformed. First, the score rewarded signals that correlate with visibility rather than with cultural relevance, so commercial and institutional accounts with strong profile metadata scored well while small collectives with sparse profiles did not. Second, the geographic component was systematically wrong in a way that was invisible until profiles were inspected by hand: several accounts scored as France-based were physically located elsewhere, and the fix required reading latitude and longitude out of the business address and testing them against a bounding box (DD-031), followed by a further narrow rule for a residual class the bounding box could not catch — Alliance Française branches in Latin American cities, whose bios legitimately mention France as a subject rather than a location (DD-032). Third, and most tellingly, an account's structural importance in the network is not the same quantity as its content volume: the network analysis found accounts ranking among the top bridging and top PageRank actors while having almost no posts loaded, their centrality coming entirely from being mentioned by others (§5.2). Any selection criterion that reads relevance off an account's own output will miss these.

Curation also had to be reconciled with discovery, since the ingestion step creates one account node per Instagram `relatedProfiles` suggestion on every scrape. Treating "not curated" as "delete" would have destroyed fresh discoveries alongside old noise. The cleanup script therefore tags, rather than deletes, any uncurated account reachable via `RELATED_TO` from a curated one, marking it `discoveredViaCuratedAccount` with a pending review status, so that these accounts survive as a review queue. The cost of the approach is stated plainly in §3.13: human curation does not scale, and the boundary of the corpus is therefore set by how much reviewing one researcher can do.

3.10 Geospatial Enrichment

Location strings extracted by the event pipeline are geocoded with Nominatim, the Open-StreetMap geocoder, chosen because it is free, requires no API key, covers both France and Colombia adequately, and imposes a rate limit (one request per second, implemented here as 1.1 seconds between requests) that is tolerable at this corpus size (DD-018). Resolved locations are placed in an administrative hierarchy — **Location** → **Arrondissement** → **City** → **Country** — with Paris arrondissements derived from postal codes in the 750XX range. Sessions are logged so that a partial run can be resumed rather than repeated.

The methodologically interesting part is what happens when geocoding fails. Roughly one event in seven never acquires coordinates, for two distinct reasons that should not be conflated: some have no location text in the caption at all, and some have text that Nominatim cannot resolve. A fallback to a city-level hint was available and was rejected as a display strategy, because it produces coordinates that look precise and are wrong — an event placed at the centre of Paris because its address could not be parsed is indistinguishable, on a map, from an event that really is there. The decision taken (§4.5) is that the discovery site shows only events with genuinely resolved coordinates. This is a presentation filter, not a deletion: the **:Event** node remains in the graph for analysis, it simply does not enter a product whose central organising principle is geography.

3.11 Ranking and Relevance Scoring

The discovery website ranks events with an explicit, documented formula rather than an opaque model, on the stated principle of contextual algorithmic curation rather than personalisation: there are no user accounts, and the same conditions produce the same ordering for every visitor. All component signals are converted to percentile ranks over the universe of publishable events before being combined. The score is

$$\text{Relevance} = 100 \cdot (0.30 Q + 0.22 A + 0.18 B + 0.20 T + 0.10 C) \cdot P$$

where Q is detection quality ($0.70 \cdot \text{eventScore} + 0.30 \cdot \text{confidence}$); A is the authority of the publishing account in the network, itself a weighted blend of PageRank (0.40), betweenness (0.25), follower count (0.20), participation coefficient (0.10) and verified status (0.05); B is social resonance (hotness score and post count); T is temporal proximity, implemented as explicit steps rather than an exponential decay so that it can be read and defended; and C is session context. P is a multiplicative penalty or bonus: political accounts $\times 0.55$, low-confidence extractions $\times 0.80$, free events $\times 1.05$.

Three aspects of this formula are methodological choices rather than tuning. First, follower count is deliberately the smallest term in A : a consulate with forty thousand followers is not thereby more culturally relevant than a collective with three thousand,

and PageRank is weighted highest precisely because it measures recognition *within the cultural network* rather than general reach. Second, Q , A , B and P are precomputed at export time because they require percentile context over the whole dataset, whereas T and C are computed in the browser on each visit because they depend on the current date and on the individual session — precomputing them would freeze the ranking to the last export (DD-043, DD-044). Third, and discussed at length in §6.2, an earlier version of the formula included a $\times 1.05$ bonus for any event carrying a non-null `culturalIdentity` value; it was removed when the project’s scope broadened. A mandatory diversification pass then caps how many events from the same author, category or parent institution can appear together at the top of the home page, without which the largest category mechanically dominates the front page.

3.12 Ethical Considerations

The corpus is restricted to public accounts and publicly visible posts; private accounts are labelled as such in the graph and are not scraped. This does not by itself settle the ethics, since public availability and consent to research use are different things. Three positions were taken in consequence. Organisational accounts — venues, institutions, collectives — are treated as public actors and named directly, since their posts are promotional by intent and naming them is what makes the mapping useful to the community. Individuals appearing incidentally in captions, comments or tags are not the object of analysis, and the project’s data-retention policy works in their favour: the cleanup phase deletes accounts never touched by manual curation together with their posts and comments, so the accumulated periphery of individual accounts discovered through mentions is removed rather than retained indefinitely.

The public artifact adds a second set of considerations, handled by design rather than by policy. The site stores no personal data on any server: it is fully static, session preferences live in the browser’s `localStorage` under a single key with no third-party cookies and no personally identifying information, and a reset control is exposed in the interface. Event cards link back to the original Instagram post rather than reproducing it, and the site presents itself as a discovery tool rather than a ticketing service. Automated detection is disclosed on each event, which matters because a language model sits in the pipeline and, as §3.13 documents, is wrong a measurable share of the time.

[TODO: Two items still to settle before submission: (i) whether the university requires a formal ethics statement or data-management plan for public-web data, and in what form; (ii) the precise status of Instagram’s terms of service with respect to scraping via a third-party actor, stated accurately rather than assumed — this is a legal-risk question the thesis should address explicitly rather than leave implied. **]**

3.13 Limitations of Method

The limitations below are drawn from failures actually observed in this project's runs, not from a generic list of caveats.

Geocoding produces false precision when it fails. Around one event in seven has no usable coordinates, and the available fallback — resolving to a city-level hint — yields plausible-looking but wrong positions. The mitigation is exclusion from the map rather than approximation (§3.10), which trades coverage for honesty.

Date extraction is the pipeline's most fragile component. Successive rounds of evaluation found dates invented from a bare month (DD-034), a theatre season "26/27" parsed as a calendar date, dates lost because the extractor truncated captions at 600 characters (DD-037), and implausible dates surviving a clamp that was silently skipped whenever either of two date parses threw — producing events dated 2036, 2052 and 2090 from text that mentioned an anniversary, an age limit and a duration in minutes (DD-039). Each was fixed once found; the pattern worth reporting is that all of them failed silently, and all were found by manual inspection rather than by any automated check.

LLM classification is inconsistent at the margin. RUN-020 recorded two nearly identical posts about the same event receiving opposite verdicts, with no identifiable code cause. Layer 3 is also capable of typing an event with a label meaning "not a cultural event" while the surrounding logic created the node anyway, because creation did not consult the category field (DD-039).

Instagram-only sourcing bounds what can be observed. Organisations that announce primarily through mailing lists, their own websites or other platforms are absent by construction, and the platform's own visibility dynamics shape which accounts are discoverable through `relatedProfiles` at all. The algorithmic layer of the network is a proprietary artefact and is analysed separately for that reason (§3.5).

Curation does not scale, and the tier labels are sparse. The trust signal that replaced automated scoring is one researcher's judgement. The consequence is visible in the network analysis: at the time of RUN-015 only 62 of 7,269 account nodes carried an explicit tier label, so tier-based comparisons describe a small hand-curated subset rather than the graph as a whole, and this must be stated wherever such comparisons appear.

Geographic scope leaks through automatic discovery. Accounts outside the intended geography enter the graph through `relatedProfiles` before curation catches them — Alliance Française branches in Colombian cities being the recurring case (DD-032) — and at least one event in the institutional sample examined in §5.5 was a restaurant menu in Spain with no connection to the project at all.

Chapter 4

System Implementation: The Hub Cultural Platform

4.1 Pipeline Architecture Overview

The system is a set of independent command-line scripts sharing a Neo4j Aura database, rather than an orchestrated application. Each script owns one phase (§3.1), reads what it needs from the graph or from disk, and writes back a disjoint set of properties. There is no scheduler and no continuous integration: runs are manual and deliberate, which suits a research setting where each execution costs money (Apify credits, LLM tokens) and where the point of a run is usually to test a change rather than to keep a service fresh.

Two properties of this architecture are load-bearing. *Incremental processing:* `extract_profiles.py` queries the graph for accounts not yet scraped rather than working from a static list, and the NLP phases skip any node whose sentinel property is already set, so a re-run costs nothing for work already done and an interrupted run can simply be restarted — which is how the connection failure mid-way through RUN-021 was recovered. *Namespace separation:* manual curation and automated ingestion write property names chosen not to collide, so re-running either never clobbers the other’s data.

4.2 Extraction and Ingestion

Extraction writes one JSON file per profile into a gitignored `data_raw/` directory; ingestion reads that directory and merges its contents into the graph. Ingestion is written entirely with `MERGE` rather than `CREATE`, and stamps every node it touches with `firstSeenAt` on creation and `lastUpdatedAt` on every write, which gives a per-node provenance trail without a separate audit table.

An instructive dead end is worth recording. Nodes created before that provenance stamping existed have neither property, and a script (`seal_legacy_batch.py`) was writ-

ten to tag them as a legacy batch so they could be identified later. By the time it ran, every node already carried some other date field, so it had nothing to seal and the `legacyBatch` property never marked a single real node. The mechanism that actually solved the underlying problem — separating material worth keeping from accumulated noise — was not a timestamp at all but the curation signal (§3.9). The lesson generalises: provenance metadata answers "when did this arrive", which turned out not to be the question that needed answering.

4.3 Network Analysis Engine

`3_analyze_network.py` exports nodes and edges to CSV, builds the two network layers in `igraph`, runs the metrics described in §3.5, and writes results back with batched `UNWIND` statements. It also supports a mode that performs the analysis alone, entirely offline from the exported CSVs, with no database connection at all. That mode is not merely a convenience: it means the analytical results of this thesis can be regenerated from files under version control, by someone with no access to the live database and no Neo4j licence, which is a stronger reproducibility guarantee than an in-database analysis could offer. Historical runs are additionally archived under `data_processed/runs/` with a timestamped directory per execution (DD-017), so that the evolution of the rankings as the corpus grew remains inspectable.

4.4 Event Enrichment Services

The event phase is where the system depends on components it does not control: two sentence-transformer models, an NLI model, and a hosted language model. Robustness therefore had to be built at the boundary. The extractor selects between a local model server and two cloud providers, marks a provider as exhausted for the remainder of a run once it has failed, and switches to the alternative rather than retrying indefinitely; a cap on rate-limit waiting prevents a provider's `Retry-After` header from stalling a run for minutes at a time (DD-033). The mechanism was validated by two distinct failure modes in practice — quota exhaustion and genuine network errors — and handled both by switching providers within ninety seconds.

4.5 The Discovery Website

The public artifact is a static site: `5_export_dashboard_data.py` reads the graph and writes a single JSON file, and `site/` contains hand-written HTML, CSS and vanilla JavaScript that consumes it, deployed on Cloudflare Pages. There is no backend, no

framework and no build step, which is a deliberate consequence of the dataset's size — the publishable events fit comfortably in a file of a few hundred kilobytes — and of the deployment constraint that the site must be reconnectable from a plain repository without configuration (DD-043, DD-044). The export is read-only and can be re-run after any extraction or resolution pass to refresh the site.

Four design decisions are worth reporting as part of the thesis's applied contribution.

The category menu is fully dynamic. There is no editorial list of families hard-coded anywhere. The "what" axis is assembled on each load from the union of the event categories and per-event artistic tags actually present in the data, with real counts, ordered by volume. The menu therefore reflects what the upstream human curation has prioritised rather than an assertion in the code, and it changes shape as the curation does.

Recommendation without accounts. Session context is stored in `localStorage`: category weights increment when a visitor filters by or opens an event in that category, decay daily, and are dropped below a floor; recently seen events are down-weighted so the home page rotates. This yields the C component of the ranking formula with no login, no server and no personal data.

Similarity without exposing embeddings. The export computes cosine similarity across all events and stores, per event, the identifiers of its five nearest neighbours. The 384-dimensional vectors themselves are discarded after use and never reach the client, which keeps the payload small and avoids publishing a representation of the corpus that was never intended to be public.

Network analysis stays under the hood. PageRank, betweenness and community structure influence the ordering of events and occasionally surface as a badge, but the first version deliberately has no dedicated organiser-profile or methodology pages. This was a scoping decision: those screens are more legible as thesis material than as product, and the remaining screens (Explore, Map, organiser profiles) are listed as future work in §7.4. **[TODO:** Include two or three screenshots here rather than only in Appendix D, once the site has been reviewed in a real browser — DD-044 records that only the pure logic of `site/app.js` has been tested in isolation, not the rendered result. **]**

Chapter 5

Results

This chapter reports results from three sources, each tied to a specific, logged run: the network analysis of RUN-015 (2026-07-24), the event extraction and resolution of RUN-021 (2026-08-13), and the dashboard export of 2026-08-14. Because the graph grows continuously, every figure below is stated with its run of origin rather than as a standing fact.

[TODO: Before submission, re-run `3_analyze_network.py run-all` and `5_export_dashboard_data.py` and refresh every number in this chapter from the resulting `data_processed/*.csv` and `site/data.json`. The figures below are correct as of the runs named, but several are known to be superseded — in particular the network analysis predates both the scope broadening and the cleanup phase. **]**

5.1 Network Structure

At RUN-015 the graph contained 7,269 `:Account` nodes. Table 5.1 reports the two layers separately, as §3.5 requires.

Table 5.1: Structure of the two network layers (RUN-015, 2026-07-24).

Metric	Social layer	Algorithmic layer
Nodes	2,973	1,706
Edges	3,348	1,956
Weakly connected components	50	6
Giant component	2,547 (85.7%)	1,537 (90.1%)
Leiden modularity ($\gamma = 0.5/1.0/1.5$)	0.902 / 0.905 / 0.905	0.860 / 0.869 / 0.868
Communities ($\gamma = 0.5/1.0/1.5$)	73 / 79 / 83	29 / 33 / 38

Both layers are dominated by a single giant component holding 85–90% of nodes, and modularity stays above 0.85 at all three resolutions tested — a stability that argues against the community structure being an artefact of one particular resolution choice.

The graph is also sparse: the social layer averages barely more than one edge per node, a property recorded as a methodological limitation from the first version of the project (DD-020) rather than discovered late.

Two cautions apply to the community counts. High modularity in a graph with many low-degree nodes can be partly mechanical, since peripheral accounts with one or two edges are assigned to small communities almost by default. The community-level E-I index gives some reassurance on this point — -0.9273 on the social layer and -0.8436 on the algorithmic layer, meaning the detected communities are strongly internally connected rather than arbitrary partitions — but does not fully resolve it.

[TODO: Run the community-size distribution check flagged in `docs/results_v2_network_analysis.md` §2 (how many of the 73–83 communities have fewer than five members) before citing the community count as a substantive finding. This is the single most important open validation in the results chapter. **]**

5.2 Bridge Actors and Community Structure

Composition first: at RUN-015 the graph was 92.5% individual accounts (6,725), with cultural institutions (182, 2.5%), state institutions (163, 2.2%), commercial accounts (151, 2.1%) and media (48, 0.7%) making up the remainder. The network is thus numerically a mass of individuals with a small institutional core — consistent with an organically formed cultural network rather than one structured top-down, though the two labelling systems involved (actor type and curatorial tier) are separate and only partially overlapping, and only 62 nodes carry an explicit tier.

Restricting to the hand-classified core, the highest participation coefficients on the social layer — the accounts whose ties span the most distinct communities — were held by an institutional communications account, a podcast, an educational-cultural account and a set of community and venue accounts. The highest PageRank scores were held by a partly different set, weighted toward dance and salsa accounts and individual artists. That the two rankings do not coincide is itself informative: being frequently referenced within a community and bridging between communities are different structural positions, and the ranking formula of §3.11 weights them separately for that reason.

The most consequential observation from this run concerns two accounts that ranked among the top bridging and top PageRank actors on both layers while having zero or one posts loaded, because their Apify scrapes returned empty. Their centrality derives entirely from being mentioned and tagged by others. This is direct empirical support for the design decision to stop inferring cultural relevance from an account’s own output (§3.9, DD-028): a content-poor account can be structurally central, and any selection rule keyed to publication volume would have discarded exactly these accounts.

[TODO: Re-validate against the current graph. The RUN-015 analysis predates both the

scope broadening and the cleanup phase, so the actor composition in particular is expected to have changed substantially — the 92.5%-individual figure counts the uncuration periphery that `cleanup_legacy_accounts.py` was subsequently designed to remove. State explicitly whether the bridging pattern survives.]

5.3 Thematic and Event Landscape

At RUN-021 the graph held 664 `:Event` nodes, of which 559 met the validity criteria for publication and 170 passed the additional requirement of resolved coordinates and entered the 2026-08-14 dashboard export alongside 200 accounts.

The distribution of the exported events by category is markedly uneven: gastronomy 51, training and workshops 27, institutional 22, visual arts 19, music 15, performing arts 9, community 9, festivals 6, audiovisual 6, politics 3 and academic 3. Gastronomy is thus the single largest category, at 30% of the exported set and 35% of the full graph.

This is a finding about the pipeline, not about the cultural field. The manual curation actively avoids gastronomy in favour of the project's declared priorities; the curator marks such accounts explicitly as "not applicable" rather than assigning them an artistic type. Gastronomy dominates because gastronomic posts are easy to scrape and easy to detect as events — they announce dates, venues and prices in exactly the register the detection layers reward — not because the project set out to collect them. The response was to demote the category in the interface rather than to filter it out of the data, keeping it available as a filter while removing it from the default prominence it would otherwise take purely on volume.

Against the four declared thematic pillars, the picture is mixed. Table 5.2 shows the distribution of curated accounts and their events by artistic theme.

Table 5.2: Curated accounts and their events by artistic theme (2026-08-13, derived from the curated `artType` field).

Theme	Accounts	Events
Visual arts	21	25
Music	23	24
Dance	21	20
Theatre	15	16
Literature	15	7
Circus	5	5
Photography	3	2
Cinema	11	2

Two of the four declared priorities — photography and cinema — have curated accounts but almost no detected events, while visual arts, music and dance, which were

never declared priorities, dominate by volume. The gap between eleven curated cinema accounts and two detected cinema events is a detection problem rather than an absence of cultural activity. A related structural gap was found in the taxonomy itself: none of the sixteen event types available to the extractor corresponds to literature, so a book launch or a poetry reading has no correct label and falls into an adjacent category or none. Since literature is one of the project’s four founding pillars, this is a pipeline gap rather than an interface detail, and it motivated the addition of per-event artistic tags (DD-042) — which, being prospective only, do not yet cover the events already in the graph.

[TODO: Re-run the theme distribution after a full extraction pass with `eventArtTags` populated, and report whether the photography/cinema gap narrows. Also decide whether the fifth, community-events category has grown enough to be reported separately here — currently it is visible only as the `comunitario` category, at 9 of 170 exported events. **]**

5.4 Geographic Distribution

Of the 559 events meeting the validity criteria on 2026-08-13, 528 (94.5%) had an associated `:Location` node and 478 (85.5%) had real coordinates. The two shortfalls have different causes: 31 events carry no location text at all, because the caption never stated one, and 50 have location text that either failed to resolve in Nominatim or had not yet been through a geocoding pass. Only the first group is irrecoverable without a change to the extraction prompt.

[TODO: Re-run `4_enrich_locations.py -dry-run` to separate the 50 unresolved locations into pending backlog and genuinely non-geocodable, then restate the coverage figure. Also report the distribution across arrondissements and communes, which the geocoding hierarchy supports but which has not yet been queried — this is the natural place to substantiate the Île-de-France scope claim of §1.4 with actual numbers. **]**

5.5 Data Quality Findings

Data quality is reported here as a result rather than deferred to a limitations footnote, because the error analysis is one of the more transferable outputs of the project.

The evaluation trajectory described in §3.8 is the headline: agreement between the pipeline and blind human classification moved from 58% (100 posts, 2026-08-09) through 82.4% precision and 85.8% recall (495 posts, 2026-08-12) to 96.7% agreement with 97.3% precision (150 posts, RUN-020). The largest single gain came from one conditional statement — making node creation depend on the language model’s own invitation and upcoming flags, which it had been computing and ignoring.

A manual review of a targeted sample sharpens the picture where the aggregate figures

are most flattering. The twenty institutional events with the highest resonance scores were classified by hand: roughly eight were genuine public events correctly categorised; roughly seven were not events at all — interviews, podcasts, press releases, appointment announcements, a private meeting — and had passed the public-invitation gate that should have stopped them; and roughly five were real events placed in the wrong category, including two visual-art installations, a concert tour, a restaurant menu in Spain unrelated to the project, and a language-course enrolment that an explicit prompt instruction was supposed to exclude. Twelve of twenty problematic is a much worse rate than the aggregate figures suggest, and the discrepancy is itself the finding: the sample was selected by highest resonance within the institutional category, which is precisely where the pipeline’s two known weaknesses concentrate — a leaky invitation gate and a tendency to type an event by the nature of the account that published it rather than by the content of the post. Aggregate accuracy over a random sample and accuracy within a targeted high-visibility slice are different quantities, and a discovery product surfaces the second.

Deduplication performance is measurable on the same run. Over 664 events, the resolver evaluated 220,116 pairs in roughly twenty seconds, identified 68 duplicate pairs, and consolidated them through 189 relationship redirections, transitive chains included. Of the pairs it rejected, 12,375 failed on similarity, 165,018 on the date window, and 6 on the geographic-conflict guardrail — the last group being precisely the failure cases that motivated the guardrail (§3.7).

[TODO: Two quantifications are still missing and would strengthen this section: (i) how many events had false-precision coordinates before the geocoding-confidence decision, obtainable by comparing a geocoding run with and without the city-hint fallback; (ii) an error rate for date extraction on the current code, since all the date figures cited above predate the DD-037 and DD-039 fixes. **]**

Chapter 6

Discussion

6.1 Interpreting Network Structure

[TODO: Depends on the re-validation flagged in §5.2. The argument to make, if the pattern holds: the network is numerically a mass of individual accounts organised around a small institutional core that does the bridging work, which suggests cultural circulation here is mediated through a handful of institutional hubs rather than through a dense peer network — with the corollary that such a topology is fragile, since removing a few hubs would fragment it. Connect to the sparse-density observation (DD-020), which reads the same structure as possible evidence of fragmentation, and be explicit about the alternative explanation: sparsity may be an artefact of how much was scraped rather than a property of the field. **]**

6.2 From a Diaspora Network to a Cultural-Relevance Network

The project’s scope changed twice. It began as a network of Colombian diaspora accounts, seeded from a single institutional account, the Colombian consulate in Paris. As the network grew through automatic relationship discovery (`RELATED_TO`), it broadened to Latin American diaspora accounts more generally. Both of these were still nationality-based criteria: an account or event qualified because of *who* was behind it.

The second, more consequential change replaced that criterion with a purely cultural one: an account or event qualifies if it is culturally relevant, with explicit priority given to photography, cinema, literature, theatre, and locally-organized community events, regardless of the nationality or cultural origin of the organizers behind it. This is a stricter methodological change than it might first appear, because it also required removing a bonus that had previously been present in the discovery website’s relevance-ranking formula (§3.11) — a multiplicative $\times 1.05$ applied to any event carrying a non-null

`culturalIdentity` value, a field which in this corpus is overwhelmingly populated with Colombian and Latin American values. Retaining that bonus after the scope change would have produced a ranking signal that was formally neutral (it ranks *cultural* events) while materially still favouring one set of national identities over every other one present in the same graph; removing it was necessary for the ranking formula to be honest about what it actually measures. The field itself was kept, as an optional filter a visitor can choose to apply — the distinction being between an identity that a user selects and one that the system applies on their behalf without telling them.

The same audit produced a negative result worth recording: community detection was confirmed to be purely topological, computed only from graph edges, so `culturalIdentity` never entered the community structure through a back door. The bias in the ranking formula was explicit and removable; the more interesting bias was not.

That one is the reference set of the event detector's first layer. Its roughly 100 phrases had been written entirely in the vocabulary of Colombian diaspora events, and the blind evaluation of 100 posts (§3.8) measured the consequence: general and French accounts lost real events at 12.9%, against 4.3% for Colombian and Latin American accounts — nearly three times the rate, on events with explicit dates and venues that were unambiguous to a human reader. A project that had already decided, at the level of stated policy, to include cultural activity regardless of origin was still, at the level of implementation, seeing that activity less well. This is the strongest argument in the thesis for treating the scope change as methodological rather than administrative: the nationality criterion was not only a boundary drawn around the corpus, it was embedded in the instrument that produced the corpus, where it was invisible until measured.

[TODO: Once network composition is re-pulled for Chapter 5, quantify here: what share of curated accounts and events carry Colombian, other Latin American, French, or other `culturalIdentity` values, before versus after the scope change. The account-level query is straightforward against the curated fields; the before/after comparison requires the archived run directories under `data_processed/runs/`. **]**

What this change gained, beyond ranking-formula honesty, is a criterion that is easier to defend and to operationalize consistently: *is this cultural, and does it fall within one of these priority domains* is a judgment that can be applied uniformly to any account the discovery process surfaces. *Is this part of the Latin American diaspora*, by contrast, required an ad hoc call for every borderline case — a French venue that regularly hosts Colombian artists, a Latin American cultural institute whose visitors are mostly French — that the pipeline had no principled way to resolve consistently. What it cost is some of the project's original sociological specificity: the initial research question was, in part, about how *one* diaspora organizes culturally in a host country, and that question is necessarily diluted once the object of study is cultural activity in general. This thesis handles that trade-off by keeping the diaspora-specific reading as one lens applied to a subset of the

network (wherever `culturalIdentity` data supports it), rather than as the network's defining boundary.

6.3 On Data Quality and the Limits of Automation

The project's most consistent lesson is that automated signals fail silently and confidently. This is not a general complaint about automation but a specific, repeated pattern in these results, visible in four independent places. The account classifier assigned high scores to accounts in the wrong country, because the geographic signal it trusted did not contain the field it assumed (DD-031). The geocoder, offered a fallback, returned coordinates that were precise and wrong rather than absent (§3.10). The date extractor filled in a day of the month that no caption ever stated (DD-034), and later accepted a theatre season as a date (DD-037). The language model typed events with a label meaning "not a cultural event" while the surrounding code created them anyway (DD-039). In none of these cases did the system report low confidence; in every case the failure was found by a human reading the output.

The corresponding lesson about human curation is more qualified than "humans are better". Human curation was the more reliable trust signal here, and replacing automated account scoring with it was the right decision on the evidence — under 10% of automatically selected accounts survived manual review. But it does not scale, it embeds a single researcher's judgement, and the evaluation protocol that validated the event pipeline measured agreement with that same judgement rather than with any external ground truth (§3.8). What the project actually demonstrates is narrower and more useful: automation was most reliable where it was constrained to deterministic, checkable operations (does this caption contain a date; do these two coordinates fall inside this bounding box; do these two country names conflict) and least reliable where it was asked to make the judgement of relevance itself. The design that emerged puts deterministic gates in front of and behind the probabilistic components, and reserves the relevance judgement for a person.

[TODO: Extend with the implication for comparable projects, stated as design guidance rather than as a moral: put a cheap deterministic gate in front of any expensive probabilistic component, prefer absence to plausible-looking approximation in any user-facing field, and instrument production runs at least as well as rehearsals (DD-038). **]**

6.4 Practical Implications: The Website as Intervention

[TODO: The discovery site is not only a dissemination channel but a sociological intervention: it makes a set of events visible together that were previously legible only to whoever followed the right accounts, and its ranking formula makes an explicit editorial claim about what counts as relevant. Two concrete hooks from the implementation are worth discussing as design research: the dynamic category menu, which delegates the taxonomy to whatever the curation prioritised rather than fixing it editorially in code; and the removal of the identity bonus, which is a worked example of an ordinary product decision that is also a political one. Connect back to the fourth objective (§1.3) — whether the artifact produced actual contact with the community. **]**

6.5 Methodological Reflections

[TODO: What would be done differently starting over. Four candidates supported by the record: (i) build the human-curation trust signal in from the start rather than arriving at it after an automated classifier underperformed; (ii) write the semantic reference set to the scope you intend rather than the scope you started with, since it silently encodes the inclusion criterion (§6.2); (iii) instrument production runs for audit from the first run, not after a run proved unauditable (DD-038); (iv) prefer explicit absence over fallback values in every field a person will read. **]**

Chapter 7

Conclusion

7.1 Summary of Contributions

[TODO: Restate the three contributions of §1.4 against what was actually delivered, with figures: the pipeline and its evaluation trajectory (methodological), the mapped network and its documented error rates (empirical), and the deployed bilingual site (applied). **]**

7.2 Answering the Research Question

[TODO: Direct answer to §1.2 and its three sub-questions, one paragraph each at most. The structural sub-question is answerable now in outline (§5.2) pending re-validation; the thematic one is answerable with §5.3 and §3.8; the practical one is answerable with §4.5 and §5.5. **]**

7.3 Limitations

[TODO: Consolidate §3.13 into boundaries on generalisation rather than a repeated list: single platform, single metropolitan area, single annotator, a corpus whose boundary is set by one researcher’s curation capacity, and a snapshot of a network that continues to change. **]**

7.4 Future Work

[TODO: From the project’s own planning notes and open items: a new account classifier with explicit geography and cultural-relevance dimensions plus anti-embeddings to penalise commercial and political content; an expanded seed set from Latin American consulates in Paris (DD-022 lists the handles, several still unverified); a revised harvest cycle

prioritising cheap profile scraping before expensive post scraping; a literature label added to the event taxonomy and a backfill decision for `eventArtTags` (DD-042); resolution of the open question in DD-035 on ongoing exhibitions; normalisation of stylised Unicode digits in captions if confirmed frequent (DD-036); and the remaining site screens — clickable *arrondissement* and commune map, Explore, Map, and organiser-profile pages, per-event bilingual text generated at extraction time, and a generic category image set.]

7.5 Closing Reflection

[TODO: Personal note tying back to the fourth objective of §1.3 — whether and how the project succeeded by its own most human metric (meeting someone through it), separate from the technical metrics.]

Appendix A

Graph Schema Reference

Table A.1: Node labels in the Hub Cultural graph.

Label	Origin	Notes
:Account	Ingestion	Also :Public or :Private; optionally :Political
:Post	Ingestion	Caption is the unit of event detection
:IgtvVideo	Ingestion	Treated like a post for authorship
:Comment	Ingestion	Author linked via WROTE
:Hashtag	Ingestion	
:Track	Ingestion	Audio track used in a post
:Location	Ingestion / events	Geocoded in phase 6
:Event	Event pipeline	Created only when the layer-3 gate passes
:Arrondissement	Geocoding	Derived from postal codes 750XX
:City	Geocoding	
:Country	Geocoding	

Table A.2: Relationship types.

Group	Types
Ingestion	PUBLISHED, HAS_HASHTAG, MENTIONS, TAGS_USER, COAUTHORED_BY, TAGGED_AT, USES_MUSIC, WROTE, ON, LOCATED_AT, RELATED_TO
NLP and geo	MENTIONS_EVENT, ORGANIZED, PARTICIPATED_IN, SUPPORTED, LOCATED_IN

[TODO: Add an entity-relationship diagram and a property-level table for :Account and :Event, distinguishing ingestion-written properties from curation-written ones (artType, culturalIdentity, geoZone, manualFollowersCount, manualDataCuratedAt) and analysis-written ones (pageRankExact, betweennessExact, participationCoef, kCore, Leiden community labels). Source: docs/data_flow.md and the COLUMN_MAP of load_manual_account_categorization.py.]

Appendix B

Selected Decision Log Entries

[TODO: Translate and reproduce in full the entries most load-bearing for the argument: DD-006 (seed strategy), DD-013 (GDS to igraph), DD-018 (Nominatim), DD-026 (seed source versus cultural target), DD-031/032 (geographic misclassification), DD-034/037/039 (date extraction failures), DD-036 (the 495-post retuning), DD-040/041 (resolver redesign and geographic guardrail), and DD-043/044 (export and site). Note in a preamble that the log is written in Spanish and that these are the author’s translations.

]

Appendix C

Evaluation Protocol and Samples

[TODO: New appendix. Reproduce the classification criterion used for the blind human labelling (including the adjudication recorded in DD-035), the per-round sample sizes and dates, and the confusion counts behind the figures in §3.8. Source files: `data_processed/eval_100_report.md`, `data_processed/claude_labels.json`, and the per-batch `eval_*.csv` files. **]**

Appendix D

Discovery Website — Screenshots

[TODO: Home screen, filter bar, event detail panel — capture once the site has been reviewed in a real browser and the outstanding punch list is cleared. **]**

Appendix E

Repository Structure

[TODO: Reproduce the phase-by-phase script table from `CLAUDE.md`'s Commands section, including the naming convention (`1_harvest_`, `2_build_`, `3_analyze_`, `4_enrich_`, `5_export_`, DD-019) and the folder conventions: `old/` for superseded or environment-incompatible scripts, `testing/` for verification and diagnostic scripts. **]**

[TODO: Create `thesis/references.bib` and uncomment the two bibliography lines above. No bibliography exists in the repository yet; the only literature reference recorded anywhere is Vertovec (2009) in DD-020, and it needs verification. **]**